

Implementing Predictive Analytics with Spark in Azure Databricks

Lab 4 – Recommenders and Clustering

Overview

In this lab, you will use Spark to build a recommender and a K-Means clustering model.

What You'll Need

To complete the labs, you will need the following:

- A web browser
- A Microsoft account
- A Microsoft Azure subscription
- A Windows, Linux, or Mac OS X computer
- Azure Storage Explorer
- An Azure Databricks cluster
- An Azure Storage account
- The lab files for this course

Note: If you have not already done so, set up the required environment for the lab by following the instructions in the [Setup](#) document for this course. Then follow the instructions in [Lab 1](#) to provision the required Azure resources.

Create a Recommender using Collaborative Filtering

First, you will evaluate a classification model that predicts whether or not a flight will be late.

Upload Source Data to Azure Storage

In this lab, you will build a recommender based on data about movie ratings. Before you can do this, you must store the data files in the storage used by your cluster. The instructions here assume you will use Azure Storage Explorer to do this, but you can use any Azure Storage tool you prefer.

1. In the folder where you extracted the lab files for this course on your local computer, in the **data** folder, verify that the **ratings.csv** and **movies.csv** files exist. These files contain data that describes 5-star rating activity from [MovieLens](#), a movie recommendation service. It was created

by GroupLens, a research group in the Department of Computer Science and Engineering at the University of Minnesota, and is used here with permission.

Note: This dataset and other GroupLens data sets are publicly available for download at <http://grouplens.org/datasets/>.

For more information, see F. Maxwell Harper and Joseph A. Konstan. 2015. [The MovieLens Datasets: History and Context](#). *ACM Transactions on Interactive Intelligent Systems (TiiS)* 5, 4, Article 19 (December 2015).

2. Start Azure Storage Explorer, and if you are not already signed in, sign into your Azure subscription.
3. Expand your storage account and the **Blob Containers** folder, and then double-click the **spark** blob container.
4. In the **Upload** drop-down list, click **Upload Files**. Then upload **ratings.csv** and **movies.csv** as block blobs to a folder named **data** in root of the **spark** container.

Create a Recommender

You will use Spark MLLib to create and evaluate your recommendation model. You can choose to work with Python or Scala.

1. In your Databricks workspace, import either **Recommender.ipynb** or **Recommender.scala**, depending on your preferred choice of language (Python or Scala), from the **Lab04** folder where you extracted the files for this course.
2. Open the notebook you uploaded and attach it to your cluster.
3. Read the notes and run the code it contains to build a classification model.

Creating a Clustering Model

Next, you will create a K-Means clustering model to segment customer data.

Upload Source Data to Azure Storage

In this exercise, you will build a clustering model based on data about customers. Before you can do this, you must store the data files in the shared storage used by your cluster.

1. In the folder where you extracted the lab files for this course on your local computer, in the **data** folder, verify that the **customers.csv** file exists. This file contains data that describes features of customers.
2. Use Azure Storage Explorer to upload **customers.csv** as a block blob to a folder named **data** in root of the **spark** container.

Create a Clustering Model

You will use a Jupyter Notebook to create your clustering model. You can choose to work with Python or Scala.

1. From the **Lab04** folder in the folder where you extracted the lab files, upload **Clustering.ipynb** or **Clustering.scala**, depending on your preferred choice of language, to your databricks workspace.
2. Open the notebook you uploaded, attach it to your cluster, and then read the notes and run the code it contains to build a clustering model.

Clean Up

Note: If you intend to proceed straight to the next lab, skip this section. Otherwise, follow the steps below to delete your Azure resources and avoid being charged for them when you are not using them.

Delete the Resource Group

1. Close the browser tab containing the databricks workspace if it is open.
2. In the Azure portal, view your **Resource groups** and select the resource group you created for your databricks workspace. This resource group contains your databricks workspace and your storage account.
3. In the blade for your resource group, click **Delete**. When prompted to confirm the deletion, enter the resource group name and click **Delete**.
4. Wait for a notification that your resource group has been deleted.
5. After a few minutes, a second resource group containing the resources for your cluster will automatically be deleted.
6. Close the browser.